


# Emotion Recognition in Still Images

## **Team Members:**

- Hanzallah Azim Burney
- Sayed Abdullah Qutb
- Balaj Saleem
- Mehmet Alper Genç
- Burak Mutlu

# Facial Recognition Dataset

- “Facial Recognition Dataset” supplied by Gotam Dahiya on Kaggle
- Emotions include happiness, sadness, anger, fear, neutral, and surprise
- Training set contains 28,079 image samples
- Testing set includes 7,178 sample images
- Each image is a 48x48 grayscale image



# ResNet50V2 and ResNet101V2

- Construct is based on pyramidal cells in the cerebral cortex [4]
- Residual building blocks allow forward and backward signals to be directly propagated from one block to any other block [2]
- Uses identity mappings as the skip connections [2]
- Solves the vanishing gradient problem that deep CNN architectures face [3]
- Reuse the activations generated by the previous layers until the adjacent layer has managed to learn its weight [4]




Fig 1. Residual Models [3]

|             | Top-1 Accuracy | Top-5 Accuracy |
|-------------|----------------|----------------|
| ResNet50V2  | 0.749          | 0.921          |
| ResNet101V2 | 0.746          | 0.928          |

Fig 2. ResNet model from Keras [1]

# InceptionV3 and InceptionResNetV2

- InceptionV3: “Greater than 78.1% accuracy on the ImageNet dataset” [5]
- Made up of numerous building blocks with different functions [5]
- Convolutions are factorized through grid size reduction and filter concatenation [6]






Figure 1. Mini-network replacing the  $5 \times 5$  convolutions.

# DenseNet Architecture

- DenseNet refers to densely connected CNNs
- A layer passes its input feature-maps to all subsequent layers; full layer-connection
- The layer at the level  $k$  has  $k$  input feature-maps
- BN - ReLU - Conv(1x1) - BN - ReLU - Conv(3x3)

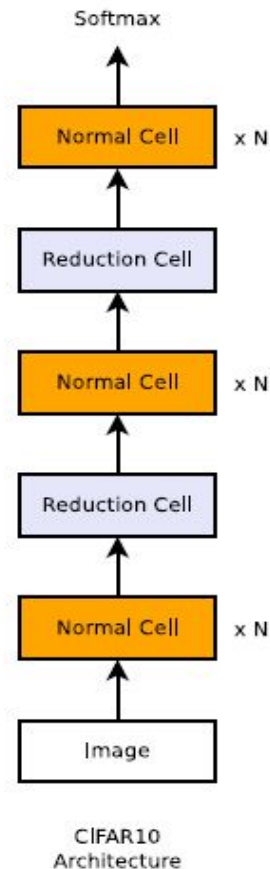
DenseNet121 - Accuracy: (0.750, 0.923)

DenseNet169 - Accuracy: (0.762, 0.932)




# NasNet

- Neural Architecture Search Network (Convolutional Neural Network)
- Large datasets, 1000+ categories
- Search for the best convolutional layer on CIFAR10, then apply on ImageNet
- Blocks found through reinforcement search method
- NasNet Large: 331x331 pixels, more parameters, approx. 11 million. 82.5%
- NasNet Mobile: Lighter, compact, mobile-platform, 224x224 pixels, less parameters approx. 5 million. 74.4%



# VGG

- State of the art Image Classification model
- Uses Convolutional Neural Networks (CNNs)
- 16-19 Layers
- Very small 3x3 convolution filters
- Pre-trained on Image-net



# Preprocessing

- 'On the fly' image preprocessing and conversion to a pixel feature set using keras
- Mean-centering, normalization, and standardization of images.
- ZCA whitening
  - Similar to dimensionality reduction using PCA
  - Retains the same spatial arrangement of pixels
  - It removes redundant pixels through a linear algebra based operation



# Evaluation Results - Epochs and Layers

- Increasing layers via using different models
- With more layers, the **model overfits** such as InceptionResNetV2 (162 layers)
- With very few layers, the **model underfits** such as VGG16 (16 layers)
- In one epoch network **sees the whole data**
- Increasing epoch allows **weight optimisation** using optimizer
- Increasing epochs can lead to **overfitting** but our models seem to avoid it

|                   | 10    | 15    |
|-------------------|-------|-------|
| ResNet50V2        | 0.477 | 0.535 |
| ResNet101V2       | 0.479 | 0.509 |
| InceptionV3       | 0.462 | 0.538 |
| InceptionResNetV2 | 0.403 | 0.514 |
| DenseNet121       | 0.502 | 0.534 |
| DenseNet169       | 0.519 | 0.516 |
| VGG16             | 0.251 | 0.251 |
| VGG19             | 0.251 | 0.251 |
| NasNetMobile      | 0.251 | 0.239 |

Accuracy Table for epochs

# Evaluation Results - Learning Rate

- Learning rate is a hyper parameter that determines the extent the weight of the model change upon each iteration.
- Following were observed:
  - Higher learning rate meant higher learning time i.e. model took longer to converge
  - Higher learning rate gave lower (or equal in rare cases) accuracies since sub optimal solution was reached
  - Keras library dynamically modifies learning rate as model trains
- Below is an example of the effect of learning rate on Inception models

| Model                 | Learning Rate | Loss  | Accuracy | Model                 | Learning Rate | Loss  | Accuracy |
|-----------------------|---------------|-------|----------|-----------------------|---------------|-------|----------|
| InceptionV3           | 1             | 1.906 | 0.200    | InceptionV3           | 0.01          | 1.205 | 0.524    |
| InceptionRes<br>NetV2 | 1             | 1.841 | 0.251    | InceptionRes<br>NetV2 | 0.01          | 1.482 | 0.403    |

# Evaluation Results - Optimizers

- Used to minimize the loss function
- **RMSprop**: exponentially weighted average of the squares of past gradients
- **SGD**: Sample batch from whole of dataset at each iteration.
- **Adam**: Combines 2 other optimizers (Momentum, RMSprop)
- RMSProp performed more accurate (due to Recurrent Neural Networks)

|                   | RMSprop | SGD   | Adam  |
|-------------------|---------|-------|-------|
| ResNet50V2        | 0.411   | 0.440 | 0.624 |
| ResNet101V2       | 0.479   | 0.466 | 0.453 |
| InceptionV3       | 0.462   | 0.472 | 0.416 |
| InceptionResNetV2 | 0.403   | 0.527 | 0.288 |
| DenseNet121       | 0.502   | 0.444 | 0.444 |
| DenseNet169       | 0.519   | 0.438 | 0.458 |
| VGG16             | 0.251   | 0.251 | 0.251 |
| VGG19             | 0.251   | 0.251 | 0.251 |
| NasNetMobile      | 0.251   | 0.177 | 0.250 |

Accuracy Table for optimizers

# Evaluation Results - Activation Functions

- Used to **rank the features** for prediction
- Activations with **smoother gradients** get better results
- Only function to get desirable results for all models is **softmax**
- **Best-performing** function for all models is softmax
- Softmax is best choice for an **image classification task**

|                   | Softmax | Relu  | Sigmoid |
|-------------------|---------|-------|---------|
| ResNet50V2        | 0.504   | 0.219 | 0.135   |
| ResNet101V2       | 0.380   | 0.343 | 0.135   |
| InceptionV3       | 0.517   | 0.426 | 0.136   |
| InceptionResNetV2 | 0.471   | 0.275 | 0.264   |
| DenseNet121       | 0.531   | 0.404 | 0.457   |
| DenseNet169       | 0.497   | 0.419 | 0.486   |

Accuracy Table for activation functions

# Evaluation Results - Normalization

- For VGG16 and InceptionV3, feature-wise normalization increased the accuracy
- ResNet50V2, ResNet101V2, InceptionResNetV2 and DenseNet models resulted in lower accuracies when data was feature-wise normalized

| Model              | Normalization | Loss  | Accuracy |
|--------------------|---------------|-------|----------|
| ResNet50V2         | True          | 1.639 | 0.411    |
| ResNet101V2        | True          | 1.346 | 0.479    |
| InceptionResNet V2 | True          | 1.482 | 0.403    |
| DenseNet121        | True          | 1.262 | 0.502    |
| DenseNet169        | True          | 1.299 | 0.519    |

| Model              | Normalization | Loss  | Accuracy |
|--------------------|---------------|-------|----------|
| ResNet50V2         | False         | 1.475 | 0.433    |
| ResNet101V2        | False         | 1.302 | 0.497    |
| InceptionResNet V2 | False         | 4.707 | 0.459    |
| DenseNet121        | False         | 1.245 | 0.525    |
| DenseNet169        | False         | 1.853 | 0.525    |

- VGG19 and NasNetMobile models had differences smaller than  $10^{-3}$  in accuracies

| Model        | Normalization | Loss  | Accuracy |
|--------------|---------------|-------|----------|
| VGG 19       | True          | 1.760 | 0.251    |
| NasNetMobile | True          | 1.176 | 0.251    |

| Model        | Normalization | Loss  | Accuracy |
|--------------|---------------|-------|----------|
| VGG 19       | False         | 1.758 | 0.251    |
| NasNetMobile | False         | 1.89  | 0.251    |

# Model Visualizations

**ResNet50V2**



**InceptionV3**




**DenseNet121**



**VGG16**



**NasNet Mobile**



# Conclusions

- Probability of random guessing was  $\sim 16.67\%$
- Deep Learning models generally performed above that
- Layers between 100 and 200 worked the best
- Epoch 15, learning rate = 0.01, SGD optimization, Softmax activation
- Emotion recognition is a viable problem to be solved using deep learning

# References

- [1] Team, K., 2020. *Keras Documentation: Keras Applications*. [online] Keras.io. Available at: <<https://keras.io/api/applications/>> [Accessed 15 November 2020].
- [2] He, K., Zhang, X., Ren, S. and Sun, J., 2016. Identity Mappings in Deep Residual Networks. *Computer Vision – ECCV 2016*, pp.630-645.
- [3] Fung, V., 2020. *An Overview Of Resnet And Its Variants*. [online] Medium. Available at: <<https://towardsdatascience.com/an-overview-of-resnet-and-its-variants-5281e2f56035>> [Accessed 15 November 2020].
- [4] En.wikipedia.org. 2020. *Residual Neural Network*. [online] Available at: <[https://en.wikipedia.org/wiki/Residual\\_neural\\_network](https://en.wikipedia.org/wiki/Residual_neural_network)> [Accessed 15 November 2020].
- [5] "Advanced Guide to Inception v3 on Cloud TPU | Google Cloud", Google Cloud, 2020. [Online]. Available: <https://cloud.google.com/tpu/docs/inception-v3-advanced>. [Accessed: 15- Nov- 2020].
- [6] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens and Z. Wojna, "Rethinking the Inception Architecture for Computer Vision", arXiv.org, 2015. [Online]. Available: <https://arxiv.org/abs/1512.00567>. [Accessed: 15- Nov- 2020].
- [7] Huang, G., Liu, Z., van der Maaten, L. and Weinberger, K., 2016. Densely Connected Convolutional Networks. [online] arXiv.org. Available at: <<https://arxiv.org/abs/1608.06993>> [Accessed 16 November 2020].